

# Lesson 12: Assessing Model Fit

Nicky Wakim

2024-05-13

# Learning Objectives

1. Use the Pearson residual statistic to determine if our preliminary final model fits the data well
2. Use the Hosmer and Lemeshow goodness-of-fit statistic to determine if our preliminary final model fits the data well
3. Use the ROC-AUC to determine how well model predicts binary outcome
4. Apply AIC and BIC as a summary measure to make additional comparisons between potential models

## Two classes ago: GLOW Study with interactions

- **Outcome variable:** any fracture in the first year of follow up (FRACTURE: 0 or 1)
- **Risk factor/variable of interest:** history of prior fracture (PRIORFRAC: 0 or 1)
- **Potential confounder or effect modifier:** age (AGE, a continuous variable)
- Fitted model with interactions:

$$\text{logit}(\hat{\pi}(\mathbf{X})) = \hat{\beta}_0 + \hat{\beta}_1 \cdot I(\text{PF} = \text{“Yes”}) + \hat{\beta}_2 \cdot \text{Age} + \hat{\beta}_3 \cdot I(\text{PF} = \text{“Yes”}) \cdot \text{Age}$$


$$\text{logit}(\hat{\pi}(\mathbf{X})) = -1.376 + 1.002 \cdot I(\text{PF} = \text{“Yes”}) + 0.063 \cdot \text{Age} - 0.057 \cdot I(\text{PF} = \text{“Yes”}) \cdot \text{Age}$$

- Let's say we are considering this model for our **final model**
- Today: determine the overall fit of this model

# Two classes ago: Reporting results of GLOW Study with interactions

- Remember our main covariate is prior fracture, so we want to focus on how age changes the relationship between prior fracture and a new fracture!

For individuals 69 years old, the estimated odds of a new fracture for individuals with prior fracture is 2.72 times the estimated odds of a new fracture for individuals with no prior fracture (95% CI: 1.70, 4.35). As seen in [Figure 1 \(a\)](#), the odds ratio of a new fracture when comparing prior fracture status decreases with age, indicating that the effect of prior fractures on new fractures decreases as individuals get older. In [Figure 1 \(b\)](#), it is evident that for both prior fracture statuses, the predicted probability of a new fracture increases as age increases. However, the predicted probability of new fracture for those without a prior fracture increases at a higher rate than that of individuals with a prior fracture. Thus, the predicted probabilities of a new fracture converge at age [insert age here].



(a) Odds ratio of fracture outcome comparing prior fracture to no prior fracture

(b) Predicted probability of fracture

Figure 1: Plots of odds ratio and predicted probability from fitted interaction model



# Assessing model fit: overview

- Once a potential final model has been determined, we need to assess the fit of the model
- Variable selection is no longer our focus at this stage
  - We want to find answer to whether the model fits the data adequately
- Assessing the Goodness of Fit or Assessing model fit
  - Assess how well our fitted logistic regression model predicts/estimates the observed outcomes
  - Comparison: fitted/estimated outcome vs. observed outcome

# Some measurements for our final model(s)

- To assess the fit of the model, it is good to have a mixture of measurements

If we want to **measure the absolute fit**: not comparing to any models, but determining if the model fits the data well

- Pearson residual statistic
  - Not the best measurement
  - Does not work if continuous variables in model
- Hosmer-Lemeshow goodness-of-fit statistic
  - Need to adjust when sample size is large
- AUC-ROC (area under the curve of the receiver operating characteristic)

If we want **comparable measures of fit**: if we want to compare candidate models that are not nested

- AUC-ROC
- AIC/BIC

# Poll Everywhere Question 1

# Components to Assess Model Fit

- The model fits the data well if
  - **Summary measurements of model fit** tell us good fit
    - Aka: Summary measures of the distance between the predicted/estimated/fitted and observed Y are small
    - Today's lecture!!
  - No major issues in **model diagnostics**, observations fit into model and our assumptions
    - The contribution of each pair (predicted and observed) to these summary measures is unsystematic and is small relative to the error structure of the model
    - Model Diagnostics that will be covered in another lecture!
- **Need to examine both components**
  - It is possible to see a “good” summary measure of the distance between predicted and observed Y with some substantial deviation from fit for a few subjects



# Summary Measures of Goodness of Fit

- Aka overall measure of fit
- What do we need to calculate them?
  - Need to calculate the fitted outcome
  - Need to calculate how close the fitted outcome is to the observed outcome
  - Summarize across all observations (or individuals' data)
- Two **tests** of goodness-of-fit
  - Pearson residual statistic
  - Hosmer-Lemeshow goodness-of-fit statistic

# Comparing *fitted outcome* to observed outcome

- In logistic regression model, we estimate  $\pi(\mathbf{X}) = P(Y = 1|\mathbf{X})$ 
  - Predicted probability,  $\hat{\pi}(\mathbf{X})$ , is between 0 and 1 for each subject
- However, we always observe  $Y = 1$  or  $Y = 0$ 
  - Not an observed  $\pi(\mathbf{X})$
- We can determine the fitted outcome by sampling  $Y$ 's from a Bernoulli distribution with the fitted probability
  - $\hat{Y} \sim \text{Bernoulli}(\hat{\pi}(\mathbf{X}))$
- If there are groups of individuals that share the *same covariate observations*, then we can use the same  $\hat{\pi}(\mathbf{X})$ 
  - $\sum_j \hat{Y} \sim \text{Binomial}(\sum_j, \hat{\pi}(\mathbf{X}))$
- Instead of comparing the expected vs. observed at individual level, we can compare them at “group” level

# Number of Covariate Patterns

- When the logistic regression model contains only categorical covariates, we can think of the number of covariate patterns
- **For example:** model contains two binary covariates (history of fracture and smoking status), there will be 4 unique combination of these factors
  - This model has 4 covariate patterns
  - Subjects can be divided into 4 groups based on the covariates' values
- We can then compute the predicted number of individuals with  $Y=1$  in each group, and compare that with the actual observed number of individuals with  $Y=1$  in that group
  - We don't need to sample this
  - We use the expected value (mean) of the Binomial to determine the  $\hat{Y}$  for each covariate pattern
  - For covariate pattern  $j$  with  $m_j$  observations:

$$\hat{Y}_j = m_j \hat{\pi}(\mathbf{X}_j) = m_j \hat{\pi}_j$$

# Learning Objectives

1. Use the Pearson residual statistic to determine if our preliminary final model fits the data well
2. Use the Hosmer and Lemeshow goodness-of-fit statistic to determine if our preliminary final model fits the data well
3. Use the ROC-AUC to determine how well model predicts binary outcome
4. Apply AIC and BIC as a summary measure to make additional comparisons between potential models



# Pearson Residual

- In logistic regression model, can use Pearson residual for summary measure of goodness-of-fit Uses the  $\hat{Y}_j$  fitted value from previous slide
- Pearson residual for jth covariate pattern is:

$$r(Y_j, \hat{\pi}_j) = \frac{(Y_j - m_j \hat{\pi}_j)}{\sqrt{m_j \hat{\pi}_j (1 - \hat{\pi}_j)}} = \frac{(Y_j - \hat{Y}_j)}{\sqrt{\hat{Y}_j (1 - \hat{\pi}_j)}}$$

- The summary statistics of Pearson residual is thus:

$$X^2 = \sum_{j=1}^J r(Y_j, \hat{\pi}_j)^2$$

# Pearson Residual procedure

1. Set the **level of significance**  $\alpha$
2. Specify the **null** (  $H_0$  ) and **alternative** (  $H_A$  ) **hypotheses**: same for all data
  - $H_0$ : model fits well
  - $H_1$ : model does not fits well
3. Calculate the **test statistic** and **p-value**
4. Write a **conclusion** to the hypothesis test
  - Do we reject or fail to reject  $H_0$ ?
  - Write a conclusion in the context of the problem

# Not going to bother going through an example

- We can calculate this by hand and test against a chi-squared distribution
- No set R code to do this
- I do not see this as the main way to determine goodness of fit... for a binary outcome!
  - Often because of the bigger issues with it... (next slide)

# Issues with Pearson Residuals

- Assume current model has  $p$  covariates...
  - then  $X^2$  (Pearson residual) follows a chi-squared distribution
    - under the null hypothesis based on large sample theory
  - **Only appropriate if the number of covariate patterns is less than the number of observations**
- When the logistic regression model contains **one or more continuous covariates**, it is likely that the **number of covariate patterns equals to the sample size  $n$**
- We should **not use Pearson Residuals** to evaluate goodness-of-fit test when the fitted **model contains one or more continuous variables**



# Learning Objectives

1. Use the Pearson residual statistic to determine if our preliminary final model fits the data well

2. Use the Hosmer and Lemeshow goodness-of-fit statistic to determine if our preliminary final model fits the data well

3. Use the ROC-AUC to determine how well model predicts binary outcome

4. Apply AIC and BIC as a summary measure to make additional comparisons between potential models

# Hosmer-Lemeshow test

- If number of covariate patterns is roughly same as the number of observations
  - Whenever you include a continuous variable in your model
  - **Hosmer-Lemeshow (HL) goodness-of-fit test** should be used instead
- However, HL test does not work well if the number of covariate patterns is small
  - HL test should not be used if the number of covariate patterns  $\leq 6$ 
    - For reference: 3 binary predictors makes 8 covariate patterns
  - Pearson residuals  $X^2$  should be used when the number of covariate patterns is small
- A large p-value from HL test suggests the model fits well

# Poll Everywhere question 2

# Hosmer-Lemeshow test

- HL test uses groupings from percentiles to basically measure what Pearson residual measures
- Steps to compute HL test statistic:
  1. Compute estimated probability  $\hat{\pi}(\mathbf{X})$  for all  $n$  subjects ( $n = 1, 2, \dots, n$ )
  2. Order  $\hat{\pi}(\mathbf{X})$  from largest to smallest values
  3. Divide ordered values into  $g$  percentile grouping (usually  $g = 10$  based on H-L's suggestion)
  4. Form table of observed and expected counts
  5. Calculate HL test statistic from table
  6. Compare HL test statistic to chi-squared distribution ( $\chi^2_{g-2}$ )



# Hosmer-Lemeshow test statistic

- The test statistic of Hosmer-Lemeshow goodness-of-fit test is denoted by  $\hat{C}$ , which is obtained by calculating the Pearson chi-squared statistic from the  $g \times 2$  table of observed and estimated expected frequencies

$$\hat{C} = \sum_{k=1}^g \frac{(o_k - n'_k \bar{\pi}_k)^2}{n'_k \bar{\pi}_k (1 - \bar{\pi}_k)}$$

- where  $n'_k$  is the total number of subjects in the  $k$ th group
- Let  $c_k$  be the number of covariate patterns in the  $k$ th decile:

$$o_k = \sum_{j=1}^{c_k} y_j$$

and

$$\bar{\pi}_k = \sum_{j=1}^{c_k} \frac{m_j \hat{\pi}_j}{n'_k}$$

# Hosmer-Lemeshow test procedure

1. Set the **level of significance**  $\alpha$
2. Specify the **null** (  $H_0$  ) and **alternative** (  $H_A$  ) **hypotheses**: same for all data
  - $H_0$ : model fits well
  - $H_1$ : model does not fits well
3. Calculate the **test statistic** and **p-value**
  - Note:  $\hat{C} \sim \chi^2_{df=g-2}$
4. Write a **conclusion** to the hypothesis test
  - Do we reject or fail to reject  $H_0$ ?
  - Write a conclusion in the context of the problem

# GLOW Study: Hosmer-Lemeshow test

- Okay, so let's look at the interaction model from last class

$$\text{logit}(\pi(\mathbf{X})) = \beta_0 + \beta_1 \cdot I(\text{PF}) + \beta_2 \cdot \text{Age} + \beta_3 \cdot I(\text{PF}) \cdot \text{Age}$$

- We need to fit the model and use a new command:

```
1 glow_m3 = glm(fracture ~ priorfrac + age_c + priorfrac*age_c,  
2               data = glow, family = binomial)  
3 library(ResourceSelection)  
4 obs_vals = as.numeric(glow$fracture) - 1  
5 fit_vals = fitted(glow_m3)  
6 hoslem.test(obs_vals, fit_vals, g = 10)
```

Hosmer and Lemeshow goodness of fit (GOF) test

data: obs\_vals, fit\_vals

X-squared = 6.778, df = 8, p-value = 0.5608

Note to Nicky: do NOT make conclusion yet! In the poll everywhere!

# Poll Everywhere question 3



# GLOW Study: Hosmer-Lemeshow test

- Conclusion: The p-value is 0.5608, so we *fail to reject* the null hypothesis that the model fits the data well. Thus, for the GLOW study data, the model with an interaction between age and prior fracture fits the data well
- Don't forget that we still need to check individual observations (Model Diagnostics!)
- R may give results for the HL test even if it is not appropriate to use it!
  - **If number of covariate patterns  $\leq 6$ , do not use HL test**

# Large sample size issue in H-L test

- When the sample size is really big ( $> 1000$ ), it is much more likely to find the H-L reject the model fit (even when the expected vs. observed in each decile categories looks pretty similar)
- This is due to “too much” power in hypothesis testing.
  - Paul et al. (2012) for samples sizes from 1000 to 25,000, the number of groups  $g$  should be equal to

$$g = \max \left( 10, \min \left\{ \frac{n_1}{2}, \frac{n - n_1}{2}, 2 + 8 \left( \frac{n}{1000} \right)^2 \right\} \right)$$

- For example, if one has a sample with  $n = 10,000$  (sample size) and  $n_1 = 1,000$  (number of events) then  $g = 500$  groups are suggested
- For  $n > 25000$ , other methods, such as partitioning data into a developmental data set (with smaller  $n$ ) and a validation set

# Final Notes on Pearson and HL tests

- They should not be used for variable selection
  - The likelihood ratio tests for significance of coefficients are much more powerful and appropriate (when nested)
- They are not for model comparison
  - One should **not** use the p-value from goodness of fit tests of different models to decide which model is better than the other
  - Something like AUC-ROC, AIC, or BIC can be used

# Learning Objectives

1. Use the Pearson residual statistic to determine if our preliminary final model fits the data well
2. Use the Hosmer and Lemeshow goodness-of-fit statistic to determine if our preliminary final model fits the data well
3. Use the ROC-AUC to determine how well model predicts binary outcome
4. Apply AIC and BIC as a summary measure to make additional comparisons between potential models



# ROC Curve and AUC (1/2)

- Receiver Operating Characteristics (ROC) curve is useful tool to quantify how good is our model predicting binary outcome
- It is a plot of sensitivity (true positive rate) versus (1-specificity) or false positive rate of fitted binary values

- True Positive Rate =  $\frac{TP}{TP + FN}$


- False Positive Rate =  $\frac{FP}{FP + TN}$

- The ROC curve shows the tradeoff between sensitivity and specificity

|                          |   | Observed outcome |    |
|--------------------------|---|------------------|----|
|                          |   | 1                | 0  |
| Predicted/fitted outcome | 1 | TP               | FP |
|                          | 0 | FN               | TN |

# ROC Curve and AUC (2/2)

- Area under the ROC curve (AUC ROC) is a reasonable summary of the overall predictive accuracy of the test
  - Accuracy means how well the predicted value matches the observed value
- The closer the curve follows the left-hand border and top border of the ROC space, the more accurate the test
  - An AUC = 1 represents 100% accuracy
- The closer the curve comes to the 45-degree diagonal line, the less accurate the test
  - An AUC = 0.5 represents an unhelpful model
    - Random predictions



# Poll Everywhere Question 4

# ROC Curve and AUC (3/3)


- Often only report the AUC
- Suggestions of how to interpret model fit through AUC values:

| AUC Values | Fit         |
|------------|-------------|
| 0.5        | Useless     |
| 0.5-0.7    | Poor        |
| 0.7-0.8    | Acceptable  |
| 0.8-0.9    | Excellent   |
| 0.9-1      | Outstanding |



# GLOW Study: ROC of interaction model

```
1 library(pROC)
2 predicted <- predict(glow_m3, glow, type="response")
3
4 # define object to plot and calculate AUC
5 rocobj <- roc(glow$fracture, predicted)
6 auc <- round(auc(glow$fracture, predicted),4)
7
8 #create ROC plot
9 ggroc(rocobj, colour = 'steelblue',
10       size = 2, legacy.axes = TRUE) +
11   ggtitle(paste0('ROC Curve ', '(AUC = ', auc, ')')) +
12   theme(text = element_text(size = 23)) +
13   xlab("False Positive Rate (1 - Specificity)") +
14   ylab("True Positive Rate (Sensitivity)")
```



- We have a poorly fitting model
- We can take `auc` and compare it to other models: good way to pick a model based on predictive power

# Another way to think about AUC

- GLOW Study: Consider the situation in which the fracture status of each individual is known
- Randomly pick one individual from fractured group and one from non-fractured outcome group
  - Based on their age and prior fracture, we will correctly predict which is from fractured group
- The AUC is the percentage of randomly drawn pairs for which we predict the pair correctly
- Therefore, AUC represents the ability of our covariates to discriminate between individuals with the outcome (fracture) and those without the outcome

# Learning Objectives

1. Use the Pearson residual statistic to determine if our preliminary final model fits the data well
2. Use the Hosmer and Lemeshow goodness-of-fit statistic to determine if our preliminary final model fits the data well
3. Use the ROC-AUC to determine how well model predicts binary outcome
4. Apply AIC and BIC as a summary measure to make additional comparisons between potential models

# AIC and BIC

- Two widely used non-hypothesis testing based measurements that helps select a good model
  - Akaike Information Criterion (AIC)
  - Bayesian Information Criterion (BIC)
- Unlike likelihood ratio test which is only suitable for nested model, **AIC and BIC are suitable for both nested and non-nested model**
- There is no hypothesis/conclusion testing for the comparison between two models
  - So not the best for selecting covariates to include in model
  - **BUT helpful if you have a few preliminary final models that you want to compare**



# Poll Everywhere Question 5

# AIC and BIC

- Both AIC and BIC penalize a model for having many parameters

| Measure of fit                       | Equation                                           | R code                       |
|--------------------------------------|----------------------------------------------------|------------------------------|
| Akaike information criterion (AIC)   | $AIC = -2 \cdot \log\text{-likelihood} + 2q$       | <code>AIC(model_name)</code> |
| Bayesian information criterion (BIC) | $BIC = -2 \cdot \log\text{-likelihood} + q\log(n)$ | <code>BIC(model_name)</code> |

- Where  $q$  is the number of parameters in the model and  $n$  is the sample size
- Both AIC and BIC can only be used to compare models fitting the same data set
- In comparing two models, the model with **smaller AIC and/or BIC is preferred**
  - When the difference in AIC between two models exceeds 3, the difference is viewed as “meaningful”

# AIC and BIC in R

- After fitting the logistic regression model, can calculate AIC and BIC
- Let's look at the AIC and BIC of our interaction model:

```
1 AIC(glow_m3)
```

```
[1] 531.2716
```

```
1 BIC(glow_m3)
```

```
[1] 548.13
```

# Example of a table for model fit statistics

```
1 model1 = glow %>% glm(formula = fracture ~ age_c + height, family = binomial)
2 model2 = glow %>% glm(formula = fracture ~ priorfrac + armassist + smoke, family = binomial)
3 model3 = glow %>% glm(formula = fracture ~ priorfrac + age_c + priorfrac*age_c, family = binomial)
4 predicted1 = predict(model1, glow, type="response")
5 predicted2 = predict(model2, glow, type="response")
6 predicted3 = predict(model3, glow, type="response")
7
8 fit_stats = data.frame(Model = c("Model 1: age + height", "Model 2: prior fracture + arm assist + smoking status", "Model 3: interaction with prior fraction + age"),
9                           AUC_ROC = c(round(auc(glow$fracture, predicted1), 4), round(auc(glow$fracture, predicted2), 4), round(auc(glow$fracture, predicted3), 4)),
10                          AIC = c(AIC(model1), AIC(model2), AIC(model3)),
11                          BIC = c(BIC(model1), BIC(model2), BIC(model3)))
12
13 colnames(fit_stats) = c("Model", "AUC-ROC", "AIC", "BIC")
14
15 fit_stats %>% gt() %>% tab_options(table.font.size = 40) %>% fmt_number(decimals = 3)
```

| Model                                                 | AUC-ROC | AIC     | BIC     |
|-------------------------------------------------------|---------|---------|---------|
| Model 1: age + height                                 | 0.642   | 542.036 | 554.680 |
| Model 2: prior fracture + arm assist + smoking status | 0.648   | 540.409 | 557.267 |
| Model 3: interaction with prior fraction + age        | 0.682   | 531.272 | 548.130 |



# Learning Objectives

1. Use the Pearson residual statistic to determine if our preliminary final model fits the data well
2. Use the Hosmer and Lemeshow goodness-of-fit statistic to determine if our preliminary final model fits the data well
3. Use the ROC-AUC to determine how well model predicts binary outcome
4. Apply AIC and BIC as a summary measure to make additional comparisons between potential models

## Summary (1/2)

| Measure of fit       | Hypothesis tested?     | Equation                                                                                       | R code                                |
|----------------------|------------------------|------------------------------------------------------------------------------------------------|---------------------------------------|
| Pearson residual     | Yes                    | $X^2 = \sum_{j=1}^J r(Y_j, \hat{\pi}_j)^2$                                                     | Not given                             |
| Hosmer-Lemeshow test | Yes                    | $\hat{C} = \sum_{k=1}^g \frac{(o_k - n'_k \bar{\pi}_k)^2}{n'_k \bar{\pi}_k (1 - \bar{\pi}_k)}$ | <code>hoslem.test()</code>            |
| AUC-ROC              | Kinda                  | Not given                                                                                      | <code>auc(observed, predicted)</code> |
| AIC                  | Only to compare models | $AIC = -2 \cdot \log\text{-likelihood} + 2q$                                                   | <code>AIC(model_name)</code>          |
| BIC                  | Only to compare models | $AIC = -2 \cdot \log\text{-likelihood} + q \log(n)$                                            | <code>BIC(model_name)</code>          |

Special notes:

- Use Hosmer-Lemshow test over Pearson residual unless number of covariate patterns is less than 6
- Cannot use Pearson residual when there is a continuous variable in the model

## Summary (2/2)

- For our interaction model:

$$\begin{aligned}\text{logit}(\hat{\pi}(\mathbf{X})) &= \hat{\beta}_0 & + \hat{\beta}_1 \cdot I(\text{PF}) & + \hat{\beta}_2 \cdot \text{Age} & + \hat{\beta}_3 \cdot I(\text{PF}) \cdot \text{Age} \\ \text{logit}(\hat{\pi}(\mathbf{X})) &= -1.376 & + 1.002 \cdot I(\text{PF}) & + 0.063 \cdot \text{Age} & - 0.057 \cdot I(\text{PF}) \cdot \text{Age}\end{aligned}$$

- We can examine the overall model fit using:
  - Not comparing to any other models:
    - Pearson residual: Not appropriate for this model
    - Hosmer-Lemeshow:  $\hat{C} = 6.778$ , p-value = 0.56
    - AUC-ROC: 0.6819
  - Can be used to compare to other models:
    - AUC-ROC: 0.6819
    - AIC: 531.27
    - BIC: 548.13

# Reference slide

- We went over model fit statistics a little in BSTA 512/612 in [Lesson 13](#)
  - Included AIC and BIC